

Outline

Video lesson 6

Task generation control

Iterative task decompositions

Recursive task decomposition

Hardware support for synchronization

Exploratory recursive problems

4/82

Outline

Video lesson 6

Task creation and synchronization (Labs summary)

Task generation control


Iterative task decompositions

Recursive task decomposition

Hardware support for synchronization

Exploratory recursive problems

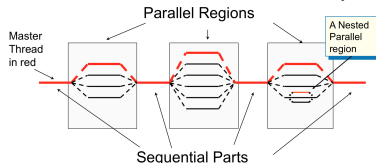
- ▶ `#pragma omp parallel`: One **implicit** task is created for each thread in the team (and immediately executed).



- ▶ `int omp_get_num_threads`: returns the number of threads in the current team. 1 if outside a parallel region
- ▶ `int omp_get_thread_num`: returns the identifier of the thread in the current team that is executing a task, a value between 0 and `omp_get_num_threads()-1`

Task creation in OpenMP (Labs summary)

- ▶ `#pragma omp parallel`: One **implicit** task is created for each thread in the team (and immediately executed). *There is a barrier synchronization at the end of the parallel region.*



- ▶ `#pragma omp task`: One **explicit** task is created, packaging code and data for (possible) deferred execution
- ▶ `#pragma omp taskloop`: **Explicit** tasks created for chunks of loop iterations. *There is a taskgroup synchronization at the end of the taskloop.*
 - ▶ In both cases, tasks executed by threads in the `parallel` region

Task ordering in OpenMP (Labs summary)

- ▶ Thread barriers: wait for all threads to finish previous work (`#pragma omp barrier` and implicit barriers at the end of OpenMP constructs)
- ▶ Task barriers:
 - ▶ `taskwait`: Suspends the execution of the current task, waiting on the completion of its **child tasks**. The `taskwait` construct is a stand-alone directive.
 - ▶ `taskgroup`: Suspends the execution of the current task at the end of structured block, waiting on the completion of **child tasks** of the current task **and their descendent** tasks.
- ▶ Task dependences (next ...)

taskwait vs. taskgroup

```
#pragma omp task {}          // T1
#pragma omp task             // T2
{
    #pragma omp task {}     // T3
}
#pragma omp task {}         // T4

#pragma omp taskwait
// Only T1, T2 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}         // T5
```

```
#pragma omp task {}           // T1
#pragma omp taskgroup
{
    #pragma omp task          // T2
    {
        #pragma omp task {}  // T3
    }
    #pragma omp task {}      // T4
}
// Only T2, T3 and T4 are guaranteed to have finished at this point when T5 is created
#pragma omp task {}          // T5
```


1. $\frac{1}{2}$ 2. $\frac{1}{2}$ 3. $\frac{1}{2}$ 4. $\frac{1}{2}$ 5. $\frac{1}{2}$ 6. $\frac{1}{2}$ 7. $\frac{1}{2}$ 8. $\frac{1}{2}$ 9. $\frac{1}{2}$ 10. $\frac{1}{2}$

Iterative task decomposition (1)

Task granularity defined by the number of iterations out of the loop each task executes. For example, using **implicit tasks**:

```
void vector_add(int *A, int *B, int *C, int n) {
    int who = omp_get_thread_num();
    int nt = omp_get_num_threads();
    int BS = n / nt;
    for (int i = who*BS; i < (who+1)*BS; i++)
        C[i] = A[i] + B[i];
}

void main() {
    ....
    #pragma omp parallel
    vector_add(a, b, c, N);
    ...
}
```

Each implicit task executes a subset of iterations, based in the thread identifier executing the implicit task and the total number of implicit tasks (i.e., number of threads in the team).


Iterative task decomposition (2) (optional)

Using the **work-sharing model** in OpenMP (not covered in this course):

```
void vector_add(int *A, int *B, int *C, int n) {
    #pragma omp for schedule(static, BS)
    for (int i=0; i< n; i++)
        C[i] = A[i] + B[i];
}

void main() {
    ....
    #pragma omp parallel
    vector_add(a, b, c, N);
    ...
}
```

Each implicit task executes chunks of iterations, depending on what is specified in the `schedule` clause. Implicit barrier at the end of each work-sharing (nowait clause to skip it).



Iterative task decomposition (3)

Task granularity defined by the number of iterations each task executes. For example, using **explicit tasks**:

```
void vector_add(int *A, int *B, int *C, int n) {
    for (int i=0; i< n; i++)
        #pragma omp task
        C[i] = A[i] + B[i];
}

void main() {
    ....
    #pragma omp parallel
    #pragma omp single
    vector_add(a, b, c, N);
    ...
}
```

each explicit task executes a single iteration of the `i` loop, large task creation overhead, very fine granularity!

Iterative task decomposition (4)

Granularity: chunk of BS loop iterations

- ▶ **Option 1:** requires loop transformation

```
void vector_add(int *A, int *B, int *C, int n) {
    int BS = ...
    for (int ii=0; ii< n; ii+=BS)
        #pragma omp task
        for (int i = ii; i < min(ii+BS, n); i++)
            C[i] = A[i] + B[i];
}

void main() {
    ...
    #pragma omp parallel
    #pragma omp single
    vector_add(a, b, c, N);
    ...
}
```

Outer loop jumps over chunks of BS iterations, inner loop traverses each chunk

Iterative task decomposition (5)

- ▶ **Option 2:** taskloop construct to specify tasks out of loop iterations:

```
void vector_add(int *A, int *B, int *C, int n) {
    int BS = ...
    #pragma omp taskloop grainsize(BS)           // or alternatively num_tasks(n/BS)
    for (int i=0; i < n; i++)
        C[i] = A[i] + B[i];
}
void main() {
    #pragma omp parallel
    #pragma omp single
    ... vector_add(a, b, c, N); ...
}
```

- ▶ `grainsize(m)`: each task executes $[min(m, n) .. 2 \times m]$ consecutive iterations, being n the total number of iterations
- ▶ `num_tasks(m)`: creates as many tasks as $min(m, n)$

List of elements, traversed using a while loop while not end of list

```
int main() {
    struct node *p;

    p = init_list(n);
    ...
    #pragma omp parallel
    #pragma omp single
    while (p != NULL) {
        #pragma omp task firstprivate(p) // see note below
        process_work(p);
        p = p->next;
    }
    ...
}
```

Note: `firstprivate` needed to capture the value of `p` at task creation time to allow its deferred execution.

Outline

Video lesson 6

Task creation and synchronization (Labs summary)

Task generation control

Iterative task decompositions

Recursive task decomposition

Reducing overheads and serialization due to synchronization

Hardware support for synchronization

Exploratory recursive problems

Recursive task decomposition: divide-and-conquer (1)

Recursively divide the problem into smaller sub-problems

```
#define N 1024
#define MIN_SIZE 64
int result = 0;


void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        result += A[i] * B[i];
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        dot_product(A, B, n);
}

void main() {
    rec_dot_product(a, b, N);
}
```


Recursive task decomposition: divide-and-conquer (2)

N=1024, MIN_SIZE=64



Recursive task decomposition: leaf strategy (1)

A task corresponds with each invocation of `dot_product` once the recursive invocations stop



- ▶ Sequential generation of tasks

Recursive task decomposition: leaf strategy (2)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)

        result += A[i] * B[i];
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}
```


Recursive task decomposition: leaf strategy (3)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    for (int i=0; i< n; i++)
        #pragma omp atomic
        result += A[i] * B[i];
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}
```

How could you reduce the overhead of updating variable `result`?

Recursive task decomposition: leaf strategy (4)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++)
        tmp += A[i] * B[i];
    #pragma omp atomic
    result += tmp;
}

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}
```

Leaf strategy: where is the task synchronization? (1)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n > MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}
```

- ▶ Where is the task synchronization?
- ▶ Are there nested tasks?

Leaf strategy: where is the task synchronization? (2)

```
#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n > MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        // Now we need the result here.
        ....
    }
}
```

- What kind of synchronization should we use? Where?

Leaf strategy: where is the task synchronization? (3)

```

#define N 1024
#define MIN_SIZE 64
int result = 0;

void dot_product(int *A, int *B, int n);


void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        rec_dot_product(A, B, n2);
        rec_dot_product(A+n2, B+n2, n-n2);
    }
    else
        #pragma omp task
        dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        #pragma omp taskwait
        // Now we need the result here.
        ....
    }
}

```

How to control task granularity in leaf strategy (1)

Leaf parallelization with **depth recursion control**



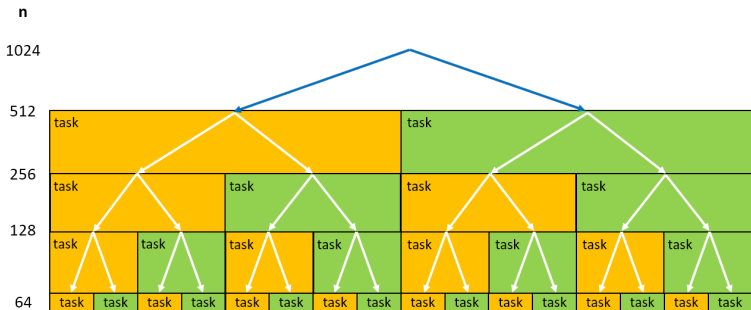
How to control task granularity in leaf strategy (2)

Leaf strategy with **depth recursion control**

```
#define CUTOFF 2
...
void rec_dot_product(int *A, int *B, int n, int depth) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth == CUTOFF)
            #pragma omp task
            {
                rec_dot_product(A, B, n2, depth+1);
                rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            }
        else {
            rec_dot_product(A, B, n2, depth+1);
            rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else // if recursion finished, need to check if task has been generated
        if (depth <= CUTOFF)
            #pragma omp task
            dot_product(A, B, n);
        else
            dot_product(A, B, n);
}
...
```

Recursive task decomposition: tree strategy (1)

A task corresponds with each invocation of `rec_dot_product`



- ▶ Parallel generation of tasks
- ▶ Granularity: some tasks simply generate new tasks

Recursive task decomposition: different sequential code ...

```
int dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    return(tmp);
}
```

```
int rec_dot_product(int *A, int *B, int n) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        tmp1 = rec_dot_product(A, B, n2);
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);
    } else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}
```

```
void main() {
    result = rec_dot_product(a, b, N);
}
```

Recursive task decomposition: tree strategy (2)

```

int dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    return(tmp);
}

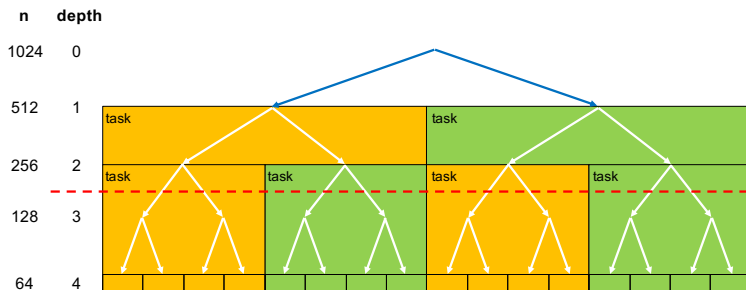
int rec_dot_product(int *A, int *B, int n) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task shared(tmp1) // firstprivate(A, B, n, n2) by default
        tmp1 = rec_dot_product(A, B, n2);
        #pragma omp task shared(tmp2) // firstprivate(A, B, n, n2) by default
        tmp2 = rec_dot_product(A+n2, B+n2, n-n2);
        #pragma omp taskwait
    } else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    result = rec_dot_product(a, b, N);
}

```

How to control task granularity in tree strategy (1)

Tree strategy with **depth recursion control**



How to control task granularity in tree strategy (2)

Tree strategy with **depth recursion control**

```
#define N 1024
#define MIN_SIZE 64
#define CUTOFF 3

int rec_dot_product(int *A, int *B, int n, int depth) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (depth < CUTOFF) {
            #pragma omp task shared(tmp1)
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            #pragma omp task shared(tmp2)
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            #pragma omp taskwait
        } else {
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else tmp = dot_product(A, B, n);
    return(tmp1+tmp2);
}
```

OpenMP support for cut-off

- ▶ `final` clause: If the expression of a `final` clause evaluates to *true* the generated task and **all of its descendent tasks** will be final. The execution of a final task is sequentially **included** in the generating task (but the task is still generated)
- ▶ `omp_in_final()` intrinsic function: it returns true when executed in a final task region; otherwise, it returns false.

OpenMP support for cut-off: tree strategy

Making use of `omp_in_final`:

```
#define MIN_SIZE 64
#define CUTOFF 3
...
int rec_dot_product(int *A, int *B, int n, int depth) {
    int tmp1, tmp2 = 0;
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        if (!omp_in_final()) {
            #pragma omp task shared(tmp1) final(depth >= CUTOFF)
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            #pragma omp task shared(tmp2) final(depth >= CUTOFF)
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
            #pragma omp taskwait
        } else {
            tmp1 = rec_dot_product(A, B, n2, depth+1);
            tmp2 = rec_dot_product(A+n2, B+n2, n-n2, depth+1);
        }
    }
    else tmp1 = dot_product(A, B, n);
    return(tmp1+tmp2);
}
...
```

Tree strategy: where is the task synchronization? (1)

Note: different tree strategy (using original seq code)... but not efficient!

```
int result = 0;
void dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i< n; i++) tmp += A[i] * B[i];
    #pragma omp atomic
    result += tmp;
}
void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task
        rec_dot_product(A, B, n2);
        #pragma omp task
        rec_dot_product(A+n2, B+n2, n-n2);
        // This taskwait is not needed: why?:
        #pragma omp taskwait
    } else dot_product(A, B, n);
}
void main() {
    #pragma omp parallel
    #pragma omp single
    rec_dot_product(a, b, N);
}
```

- ▶ Where is the task synchronization?
- ▶ Are there nested tasks?

Tree strategy: where is the task synchronization? (2)

```
int result = 0;
void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task
        rec_dot_product(A, B, n2);
        #pragma omp task
        rec_dot_product(A+n2, B+n2, n-n2);
    } else dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        rec_dot_product(a, b, N);
        // Now we need the result here.
        ....
    }
}
```

- What kind of synchronization should we use? Where?

Tree strategy: where is the task synchronization? (3)

```

int result = 0;
void dot_product(int *A, int *B, int n);

void rec_dot_product(int *A, int *B, int n) {
    if (n>MIN_SIZE) {
        int n2 = n / 2;
        #pragma omp task
        rec_dot_product(A, B, n2);
        #pragma omp task
        rec_dot_product(A+n2, B+n2, n-n2);
    } else dot_product(A, B, n);
}

void main() {
    #pragma omp parallel
    #pragma omp single
    {
        #pragma omp taskgroup
        {
            rec_dot_product(a, b, N);
        }
        // Now we need the result here..
        ....
    }
}

```


Avoiding task barriers: task dependences (1)

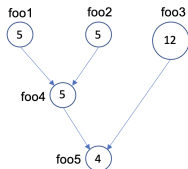
- ▶ The OpenMP runtime detects dependences between sibling tasks (i.e. from the same parent task) through the specification of the directionality for the variables used in the tasks

```
#pragma omp task [depend (in : var_list)]
                  [depend (out : var_list)]
                  [depend (inout : var_list)]
```

Task dependences are derived from the directionality type (`in`, `out` or `inout`) and its items in `var_list`; this list may include array sections (e.g. `v[0:n]`)

Serialisation caused by task barriers (1)

Given a TDG to implement with the OpenMP tasking model:

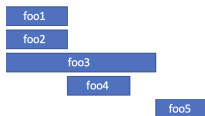


```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp task
foo3()
#pragma omp taskwait
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```

```

#pragma omp task
foo1()
#pragma omp task
foo2()
#pragma omp taskwait
#pragma omp task
foo3()
#pragma omp task
foo4()
#pragma omp taskwait
#pragma omp task
foo5()
  
```



- ▶ `taskwait` with `depend` clause: instead of waiting for all child tasks to complete execution, it only waits for the predecessor child tasks according to the `in`, `out` and `inout` specifiers

```
int x=0; y=2;
```

```
#pragma omp task depend(out: x) shared(x)
compute_short1(&x);
```

```
#pragma omp task shared(y)
compute_long(&y);
```

```
#pragma omp taskwait depend(in: x)    // y not waited for at this point
printf("intermediate value for x=%d\n",x);
```

```
#pragma omp task shared(x)
compute_short2(&x)
```

```
#pragma omp taskwait
printf("final values for x=%d ; y=%d\n", x, y);
```


- ▶ An iterator can be used in the depend clause, expanding to multiple values in the specifier they appear

```

for (i = 0; i < n; ++i)
    if (i%2) {
        #pragma omp task depend(out: v[i])
        compute_element(&v[i], i);
    }

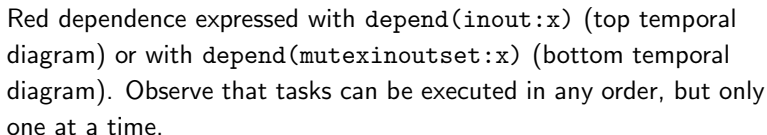
#pragma omp task depend(iterator(it = 0:n), in: v[it])
// could also be depend(iterator(it = 1:n:2), in: v[it])
odd = sum_odd_elements(v, n);

even = sum_even_elements(v, n);

```

Note: this is not equivalent to the use of an array section in the `in` specifier (i.e. `depend(in:v[0:n])`), why not?

- ▶ `mutexinoutset` specifier: equivalent to `inout` but all dependent tasks can be executed in any order, one after the other



Cancellation points in OpenMP: very simple example (optional)

```
#pragma omp taskgroup
for (i=0; i<1000; i=i+100)
    #pragma omp task firstprivate(i) private(j)
    {
        for (j=i; j<i+100; j++) {
            if (do_computation(j) == 0) {
                #pragma omp cancel taskgroup
            }
            #pragma omp cancellation point taskgroup
        }
    }
```

The first task with 0 as a result of `do_computation` will finalise the execution of all the tasks in the `taskgroup`

Protecting task interactions in OpenMP (Labs summary)

Two mechanisms:

1. Atomic accesses: mechanism to guarantee atomicity in load/store instructions

```
#pragma omp atomic [update | read | write]
    expression
```

- ▶ Atomic updates: `x += 1`, `x = x - foo()`, `x[index[i]]++`
- ▶ Atomic reads: `value = *p`
- ▶ Atomic writes: `*p = value`

Reducing task interactions: overhead (1)

Reductions: replicate key data structures and locally working with these local structures; when appropriate, locally replicated data structures are combined into the final global result

```
int result = 0;
// Assume this function is instantiated as a task
void dot_product(int *A, int *B, int n) {
    for (int i=0; i < n; i++)
        #pragma omp atomic
        result += A[i] * B[i];
}
```

could be easily transformed into

```
void dot_product(int *A, int *B, int n) {
    int tmp = 0;
    for (int i=0; i<n; i++)
        tmp += A[i] * B[i];

    #pragma omp atomic
    result += tmp;
}
```

Reducing task interactions: overhead (2)

Specifying reduction operations in explicit tasks generated with either `task`:

```
#pragma omp parallel
#pragma omp single
{
    #pragma omp taskgroup task_reduction(+: sum)
    for (i=0; i< SIZE; i++)
        #pragma omp task firstprivate(i) in_reduction(+: sum)
        sum += X[i];
}
```

or taskloop:

```
#pragma omp parallel
#pragma omp single
{
    // implicit taskgroup in taskloop construct
    #pragma omp taskloop reduction(+: sum)
    for (i=0; i< SIZE; i++)
        sum += X[i];
}
```


Low-level synchronization functions using *locks*

Locks: special variables that live in memory with two basic operations:

- ▶ Acquire: while a thread has the lock, nobody else gets it; this allows the thread to do its work in private, not bothered by other threads
- ▶ Release: allow other threads to acquire the lock and do their work (one at a time) in private

Type definition and intrinsics:

```
void omp_init_lock(omp_lock_t *lock)
void omp_destroy_lock(omp_lock_t *lock)

void omp_set_lock(omp_lock_t *lock)
void omp_unset_lock(omp_lock_t *lock)

int omp_test_lock(omp_lock_t *lock)
```


Easily parallelizable using an iterative task decomposition using `taskloop`. However ...

- ```
typedef struct {
 int data;
 element *next;
} element;

int dataTable[SIZE_TABLE];
element * HashTable[SIZE_HASH];

#pragma omp taskloop
for (i = 0; i < elements; i++) {
 int index = hash_function (dataTable[i], SIZE_HASH);
 #pragma omp critical // atomic not possible here
 insert_element (dataTable[i], index, HashTable);
}
```

- 59/82















```

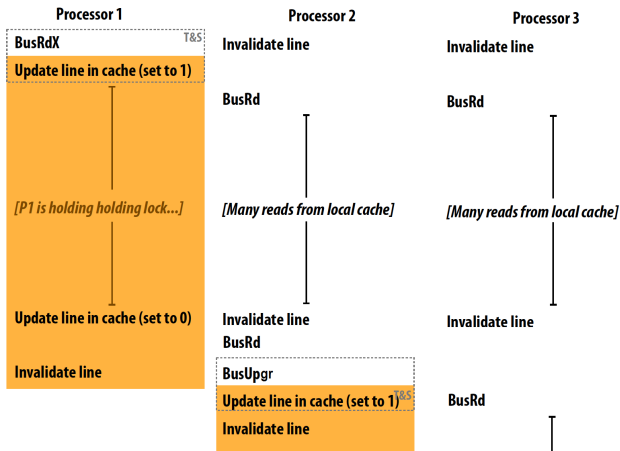
 set_lock: t&s r2, flag // test and acquire lock if free
 bnez r2, set_lock // do it again if already locked
 ...
unset_lock: st flag, #0 // free the lock

```





## test-test-and-set lock coherence traffic



## Support for synchronization at the architecture level

- ▶ Atomicity difficult or inefficient in large systems. Alternative:  
**Load-linked Store-conditional ll-sc**
  - ▶ ll returns the current value of a memory location
  - ▶ sc stores a new value in that memory location if no updates have occurred to it since the ll; otherwise, the store fails
  - ▶ sc returns success (1) or failure (0)
- ▶ Examples implementing atomic exchange (left) and fetch-and-increment (right):

```
// exchange r4 with location.
try: mov r3, r4
 ll r2, location
 sc r3, location
 beqz r3, try
 mov r4, r2
```

```
// add 1 to location
try: ll r2, location
 add r3, r2, #1
 sc r3, location
 beqz r3, try
```

- ▶ test-test-and-set technique can also be implemented with `ll-sc`
  - ▶ First, wait using load linked instruction `ll` (lock will be cached)
  - ▶ Second, use store conditional `sc` operation to test if someone else did it first

- ▶ How to implement a barrier synchronization primitive?
  - ▶ Threads arriving wait until all have reached the barrier
  - ▶ Structure with fields {lock, counter, flag}

Does it work when consecutive barriers appear? Try to solve it





# How would you address the N-queens problem? (1)

```


char *a; // Solution being explored
int sol_count = 0; // Total number of solutions found
int size = 8; // board size

void nqueens(int n, int j, char *a) {
 if (j == n) sol_count += 1;
 else
 // try each possible position for queen <j>
 for (int i=0 ; i < n ; i++) {
 a[j] = (char) i;
 if (ok(j + 1, a)) nqueens(n, j + 1, a);
 }
}

int main() {
 a = alloca(size * sizeof(char));
 nqueens(size, 0, a);
}


```

a = [0, 6, 3, 5, 7, 1, 4, 2]



## How would you address the N-queens problem? (2)

For a 4x4 board, the recursion tree would be ...



```

void nqueens(int n, int j, char *a) {
 if (j == n)
 #pragma omp atomic
 sol_count += 1;
 else
 // try each possible position for queen <j>
 for (int i=0 ; i < n ; i++) {
 a[j] = (char) i;
 if (ok(j + 1, a))
 #pragma omp task // all firstprivate by default
 nqueens(n, j + 1, a);
 }
 // Do we need to insert a task barrier at this point?
}

int main() {
 a = alloca(size * sizeof(char));
 #pragma omp parallel
 #pragma omp single
 nqueens(size, 0, a);
}

```

Do we need a new board for each task to be able to explore its own path? Is the implicit `firstprivate(a)` enough?






```

void nqueens(int n, int j, char *a) {
 if (j == n)
 #pragma omp atomic
 sol_count += 1;
 else
 // try each possible position for queen <j>
 if (!omp_in_final()) {
 for (int i=0 ; i < n ; i++) {
 a[j] = (char) i;
 if (ok(j + 1, a))
 // allocate a temporary array and copy <a> into it
 char * b = alloca(n * sizeof(char));
 memcpy(b, a, (j + 1) * sizeof(char));
 #pragma omp task final(j>CUT_OFF)
 nqueens(n, j + 1, b);
 }
 #pragma omp taskwait
 } else
 for (int i=0 ; i < n ; i++) {
 a[j] = (char) i;
 if (ok(j + 1, a)) nqueens(n, j + 1, a);
 }
}

```

## Another example: 15-puzzle (without code) ... (optional)

The solution to a 15-puzzle (a tile puzzle). Possible movements of the empty cell: UP, RIGHT, LEFT and DOWN. Here we show a series of moves that transform a given initial state to the desired final state:








## Another example: 15-puzzle (without code) ... (optional)

Anomalous speed-ups of the parallel formulation of the problem:  
the speed-up depends on where the solution is found ...



- ▶ Left:  $T_1 = 2 \times m + 1$  and  $T_4 = 1$ , therefore ...  $S_4 = 2 \times m + 1$
- ▶ Right:  $T_1 = m$  and  $T_4 = m$ , therefore ...  $S_4 = 1$

And the parallel efficiency (i.e. how well used are processors)?  
Observe that on the right three processors waste their computation

## Parallelism (PAR)

Mastering your task decomposition strategies:  
going some steps further

Eduard Ayguadé, José Ramón Herrero (Q1),  
Daniel Jiménez and Gladys Utrera

Computer Architecture Department  
Universitat Politècnica de Catalunya

Course 2022/23 (Spring semester)